

Claude Certified Associate – Foundations: Exam Cheatsheet

Companion to the 2-hour O'Reilly exam-prep course. Exam facts follow the official Exam Guide v1.0 (effective July 2026) and the Partner Academy FAQ and policies.

60 items

Multiple-choice and multiple-response. Each item states how many to pick.

120 min

About 2 min per item. Plan for ~135 min of seat time.

720 / 1,000

Scaled cut score on a 100–1,000 scale. A scaled 720 ≠ 72% raw.

\$99 USD

Per attempt. Partner-tier discounts may apply.

12 months

Valid 12 months. On-time renewal is free and non-proctored.

DELIVERY

Pearson VUE, online proctored or at a test center.

RULES

Closed book. English only. No AI tools, notes, or translation tools.

SCORE REPORT

Pass/fail, scaled score, and percent correct by domain. Only the total decides.

RETAKES

Wait 14 / 30 / 90 days after fails 1 / 2 / 3. Max 4 attempts per rolling 12 months.

RESCHEDULE

Change or cancel ≥ 48 h ahead. The guide says 24 h; the FAQ and policies say 48 h.

Eligibility. Registration currently requires a Claude Partner Network company email. Personal addresses do not work. The skills on this sheet still apply to anyone who uses Claude at work.

BLUEPRINT: 7 DOMAINS, 30 OBJECTIVES

#	Domain	Weight	~Items	Lesson
D1	Prompting and Task Execution	14%	8	02
D2	Output Evaluation and Validation	21%	13	03
D3	Product and Model Selection	12%	7	01
D4	Workflow Integration and Solution Design	16%	10	04
D5	Configuration and Knowledge Management	12%	7	05
D6	Governance, Risk, and Responsible Use	15%	9	06
D7	Troubleshooting and Optimization	10%	6	07

Item counts are weight × 60, rounded. The guide publishes weights only. **D2 alone is 21%.** D1 + D2 = 35%. D2 + D4 + D6 = 52%: judgment about accuracy, workflow fit, and risk outweighs feature trivia.

CORE EXAM LOGIC

- 1 Verify specifics against the source.** Model confidence is not evidence.
- 2 Match the model and the feature to the task.** Faster tier for volume, top tier for hard reasoning.
- 3 Anonymize before upload.** "Tell Claude not to retain it" is not a control.
- 4 Put recurring context in a Project.**
- 5 Diagnose what changed before rewriting.** Change one thing at a time.
- 6 Escalate** integrations to Developers or Architects, and high-stakes judgments to qualified humans.
- 7 Pick the proportionate step** that lets the work continue safely.

How items are written: a workplace scenario with a role → a constraint in the stem (audience, cost, volume, data policy) → "most appropriate" or "best fits" → four options. One option is proportionate. Three fail in predictable ways. **"Most appropriate" does not mean "most thorough."**

THE FIVE DISTRACTOR PATTERNS: NAME THE PATTERN, CROSS IT OUT

Pattern	What it sounds like	Why it fails
Blind trust	"Send it as-is, Claude was confident." "It cited a source." "It's internal."	Confidence, citations, and polish are not evidence.
Fake safeguard	"Ask Claude to rate its confidence." "Tell Claude not to retain it." "Use Incognito." "Add a disclaimer."	It looks like a control and changes nothing about the risk.
Irrelevant fix	"Reword it more formally." "Switch AI platforms." "Put the rule in the Project description."	It solves a problem the stem does not have.
Over-reaction	"Skip the analysis." "Ban AI for the team." "Rebuild from scratch."	A proportionate path exists.
Over-engineering	"Top model for every reply." "Escalate everything to Legal." "Fully automate on day one."	It ignores cost, speed, or the stated scope.

THE AI FLUENCY 4DS AS A SECOND LENS

4D	Domains	The question wants...
Delegation	D3, D4	a split of work between people and Claude
Description	D1, D5	clearer context, instructions, setup
Discernment	D2, D7	verification or diagnosis
Diligence	D6, D2	a real control or a disclosure

CUE WORDS IN THE STEM → FIRST INSTINCT

"every week", "same format", "the team"	Project
"customer", "regulator", "legal", "executive"	Verify, then qualified human review
"names", "accounts", "PII", "policy restricts"	Anonymize first
"last month", "current", "latest"	Web search or Research, check citations
"hundreds", "routine", "fast", "cheap"	Faster, lower-cost model
"integrate", "automate with our system"	Escalate to Developer or Architect
"it worked until..."	Find what changed

OBJECTIVES (VERBATIM)

1. Create effective prompts for business and technical tasks
2. Apply task decomposition techniques to structure complex requests

3. Iterate prompts to improve output quality
4. Adapt prompting strategies based on task type (analysis, research, drafting, brainstorming)

MENTAL MODEL

Treat Claude as "a brilliant but new employee who lacks context on your norms and workflows." Golden rule: if a colleague with minimal context would be confused by your prompt, Claude will be too. Weak output on the exam usually traces to a prompt a new colleague could not act on.

PROMPT ANATOMY

Role	"You are an operations analyst."
Context	Who reads it, why it matters: "VP reads this on mobile before Monday."
Task	"List every action item with owner and due date."
Inputs	"Use only the attached notes."
Constraints	"Under 120 words. No jargon."
Format	"Table: Item, Owner, Due, Status."
Examples	One filled-in example row or paragraph

Add how to collaborate: "If it isn't in the notes, write 'not stated'." "Push back if my assumptions are wrong." "Tell me what you're uncertain about."

FOUR PRINCIPLES

- **Concrete limits beat adjectives.** "Under 100 words" works. "Be concise" barely moves output.
- **An example is the strongest format signal.** Show a filled template.
- **Explain why.** A reason generalizes better than "NEVER do X".
- **State the goal and the format.** "Three bullets" is a format. "Convince my team the timeline is realistic" is a goal.

Long prompts: separate parts with labels or XML-style tags (<context>, <notes>, <task>, <format>). Long documents: material first, question at the end.

STRATEGY BY TASK TYPE

Type	Include	Avoid
Analysis	Data, criteria, reasoning before conclusions, ask for one recommendation	"What do you think?" with no criteria
Research	Research mode or web search, goals, sections, constraints, checkable sources	Trusting training knowledge for recent facts
Drafting	Audience, purpose, tone, length, format, an example	Generic "write a post about X"
Brainstorming	Volume first (20 ideas), then narrow with criteria, turn by turn	"The best idea" in one shot

DECOMPOSITION AND PROMPT CHAINING

Sequential steps	"First A, then B, then C." Simple flows you trust.
Checkpoint	"Stop after step 2 and show me." Later steps depend on a result you must verify.
Prompt chain	Separate prompts, each using the checked output of the last. Multi-part deliverables.
Self-correction	Draft → review against stated criteria → revise.

Decompose when there are several deliverables, dependent steps, or a needed mid-point check. Each step should produce something you can check.

ITERATION

Give **specific** feedback. "Make it shorter" → "Cut the first two paragraphs and end with the ask."

- **Follow up** in the same chat when the draft is mostly right.
- **Edit the original prompt** when the first prompt was the problem.
- **Start fresh** when the chat has gone off track.

Stuck? "Ask me the questions you need answered to write a better prompt." Instruction followed literally but uselessly? Restate the goal.

D1 TRAPS

- "Be detailed", "be concise", CAPITALS: intensifiers add no structure.
- Regenerate until it works: the cause stays in place.
- One giant prompt for a multi-part deliverable: decompose.
- Paste every document "for context": attention is finite, so curate.
- Bigger model to fix generic output: add context and constraints.
- Format described in prose when it must match a template: show an example.
- Repeat an ignored instruction louder: restate the goal.

OBJECTIVES (VERBATIM)

1. Evaluate Claude-generated outputs for accuracy and completeness
2. Identify hallucinations, inconsistencies, and biases in responses
3. Apply fact-checking and validation techniques

4. Determine when human review or additional verification is required
5. Edit, adapt, refine, and compare outputs for the intended audience
6. Organize and curate information and select appropriate output formats (artifacts, inline, structured data)

ACCURACY AND COMPLETENESS ARE SEPARATE CHECKS

Accuracy: is every statement true per the source of record?

Completeness: did the output include everything that matters?

An output can be 100% accurate and still fail by omission. You catch omissions only by reading the source.

"Fabrication concentrates in specificity: names, dates, statistics, citations, URLs, quotes." The more precise a claim, the more it needs checking.

FAILURE TYPES

Hallucination	Specific-looking detail with no basis: cites "Section 8.2" of a rule with seven sections
Distortion	A real fact subtly changed: "up to \$50,000" → "start at \$50,000"
Omission	A material point left out: a legal-hold exemption disappears
Inconsistency	Contradictions inside one answer, or different answers across runs
Bias	One-sided framing, stereotyped defaults, unsupported group generalizations
Sycophancy	Agrees with a wrong premise or caves under pushback
Stale knowledge	Out-of-date training facts presented as current

SIGNALS THAT ARE NOT EVIDENCE

- Claude sounds confident, or rates its own confidence high (official Sample 1 rejects this).
- Claude re-checks its own answer: the check is not independent.
- Two runs agree: consistency ≠ accuracy. Two runs that **disagree** are a useful warning.
- It cites a source: you still have to open it.
- Polished, formatted, or in an artifact: people check polished output less often.

ADAPT FOR THE AUDIENCE

- **Executive:** bottom line first, decision needed, two or three numbers.
- **Practitioner:** steps, owners, details, exceptions.
- **Customer or public:** plain language, no jargon, reviewed before sending.

Compare drafts against the goal and the reader. Length and polish are weak signals.

FACT-CHECKING TECHNIQUES

Quote first	"For each claim, quote the exact supporting sentence."
Cite location	"Give the section number for every requirement."
Permission not to know	"If the source doesn't say, write 'not in source'."
Restrict sources	"Use only the attached documents."
Label unsourced	"Mark general-knowledge claims 'unsourced'."
Compare runs	Same prompt in fresh chats. Disagreement flags weak claims.
Ground current facts	Web search or Research, then open the cited pages.
Compute exactly	Code execution for arithmetic. You still check inputs.
Source of record	A person reads the original. This is the verification.

OUTPUT-FORMAT PICKER

Inline reply	Quick answer read once: "Which section covers access logs?"
Artifact	Standalone deliverable to edit, reuse, or share: brief, tracker, dashboard
File creation	Downloadable .docx, .xlsx, .pptx, or PDF
Structured data	Feeds another step or spreadsheet: table or CSV with fixed columns

Ask for the deliverable. "Summarize Q3" gets a chat reply. "Turn Q3 into a one-page doc for leadership" gets an artifact.

HUMAN-REVIEW TRIGGERS: SCALE SCRUTINY TO STAKES

Output	Minimum check
Brainstorm, draft for yourself	Skim for usefulness
Internal summary others act on	Verify names, numbers, dates, quotes
External (customers, press, regulators)	Verify specifics + accountable owner reviews before it goes out
High-risk domains*	Qualified professional reviews + disclose AI involvement
Decisions about people	Claude organizes information. A human owns the decision.

*Usage Policy: legal, healthcare, insurance, finance and lending, employment and housing, academic testing, media. A reviewer counts only if qualified, sees the evidence, can stop the work, and reviews before the consequential step.

D2 TRAPS

- Send it because it was confident or cited a section *blind trust*
- Ask Claude to rate or double-check itself *fake safeguard*
- Reformat, reword, or add a disclaimer *irrelevant fix*
- Verify one claim, assume the rest *blind trust*
- Fill a "data unavailable" gap with an estimate *blind trust*
- Throw out the whole output for one error *over-reaction*
- Legal review of an internal brainstorm *over-engineering*

OBJECTIVES (VERBATIM)

1. Select appropriate Claude product features (Projects, research mode, chat, artifacts)
2. Differentiate between Claude model types (Haiku, Sonnet, Opus)
3. Align model selection with task requirements (cost, speed, quality)
4. Understand and manage context limitations and memory considerations (when to restart, summarize, or persist)

FEATURE PICKER

You need to...	Use	Signal words
Think, draft, brainstorm turn by turn	Chat	"explore", "quick draft", "talk through"
Reuse the same files and instructions across chats	Project	"every week", "same format", "team", "ongoing"
Get one current fact fast	Web search	"today's", "latest price", "who won"
Produce a cited report from many sources	Research (web search on)	"comprehensive", "compare vendors", "market scan", "with sources"
Reason hard over material you have	Extended thinking / higher effort	"complex logic", "trade-offs", "multi-step analysis"
A standalone deliverable to edit, reuse, share	Artifact	"one-pager", "dashboard", "tracker", "share with the team"
A downloadable Word, Excel, PowerPoint, PDF	File creation	".docx", "spreadsheet", "slide deck to send" (needs code execution and file creation on)
Multi-step task ending in a finished file, in your folders, or on a schedule	Cowork (desktop, paid plans)	"every Monday", "in this folder", "finished file"
Write, test, ship software	Claude Code	"code", "repository", "deploy": developer territory, so the Associate escalates

Research is for depth; a quick fact needs web search. **Projects are for recurring context.** **Features combine:** a weekly competitor brief can use a Project + Research + an Artifact.

MODEL PICKER: LEARN THE TIERS AND THEIR TRADE-OFFS

Tier	Speed / cost	Best for	Exam cue
Haiku	Fastest / lowest	Quick answers, simple extraction, high-volume routine drafts	"hundreds of", "simple", "fast"
Sonnet	Fast / moderate	Everyday writing, analysis, multi-step work. The default.	"typical", "most tasks"
Opus	Slower / highest	Deep reasoning, complex analysis, hard judgment calls	"complex", "nuanced"

Sample 2 logic: faster, lower-cost model for straightforward high-volume work; the most capable model for complex reasoning. Start with Sonnet and move to Opus on evidence. Anthropic's lineup also includes **Fable**, a newer tier above Opus. The exam guide names three tiers, so answer with Haiku, Sonnet, Opus.

THE EFFORT LEVER

Model	Changes capability and cost
Effort / extended thinking	Changes how thoroughly Claude works before answering. Cost rises. Core capability stays the same.

Too little effort: missed instructions, unfinished work. Too much: verbose, scope creep, slow simple replies. Think in **cost per task**: one good turn can beat five cheap follow-ups. A bigger model does not fix a missing source, a vague prompt, or a policy problem.

CONTEXT WINDOW FACTS

- Holds your messages, Claude's replies, and uploads. Standard: roughly 500 pages of text.
- **The whole conversation is re-sent every turn.** Long chats cost more and slow down.
- **Attention is finite.** Facts buried mid-material get missed. Key instructions go at start and end.
- **Long chats drift.** Early rules weaken.
- Near the limit, claude.ai summarizes earlier messages. Summaries lose detail.
- **Context limit** = what fits in one chat. **Usage limit** = how much you can use over time.

RESTART, SUMMARIZE, OR PERSIST

Chat drifted, ignores early rules	Restart with a short summary of decisions and rules
Same task, chat getting long	Summarize: handoff summary, check it, carry into a new chat
New, unrelated task	Restart: fresh chat, no old bias
Facts or rules needed across sessions	Persist: Project instructions/knowledge, personal instructions, memory
Sensitive one-off, keep out of history	Incognito chat

MEMORY AND INCOGNITO

- Claude does not learn from your corrections. A correction lasts only if persisted.
- **Memory** is on by default on Free, Pro, Max. Team/Enterprise owners control it. Each Project has its own memory. View, edit, delete in Settings.
- **Search and reference chats** (paid plans) looks up past conversations.
- **Incognito:** not saved to history, not searched later, still follows org retention. It does not make restricted data safe to upload.
- Memory is a convenience. It does not replace Project knowledge as a source of truth.

D3 TRAPS

- Most capable model for everything *over-engineering*
- Upgrade the model when facts or context are missing *irrelevant fix*
- Research mode for one quick fact *over-engineering*
- "Claude will remember last week": only via memory, a Project, or instructions
- Repeat the format rule louder in a 60-turn chat: restart and persist
- "Incognito means the data is approved" *fake safeguard*
- One long chat for everything: costs more, drifts

OBJECTIVES (VERBATIM)

1. Apply Claude to analyze requirements and use cases
2. Leverage Claude for research, planning, and process optimization

3. Use Claude to support solution design, development, and iteration
4. Integrate Claude into existing workflows to augment or redesign them
5. Communicate Claude's value and limitations to stakeholders

REQUIREMENTS FIRST

Deliverable	What exactly gets produced?
Audience	Who uses it, for what decision?
Sources	What data goes in, and is it approved?
Success	How will you know it works?

Audit: how often, how long, how standardized, cost of imperfection?
Have Claude interview you first.

WHERE CLAUDE FITS

Strong fit	Weak fit
Drafting from approved inputs	Final authority on facts
Summarize, categorize, extract	Accountable decisions (hire, escalate an account)
Repurpose for formats and audiences	Emotionally sensitive conversations
Brainstorming, first-pass plans	Precise high-volume math a spreadsheet should own
Research synthesis with citations	Live data without search, Research, or a connector

THE DELEGATION MAP: THREE BUCKETS

Claude handles	Standardized, documented, repeatable, low cost of error. Spot-check.
Claude assists, human decides	Claude drafts. A person reviews before it counts.
Human handles	High stakes, sensitive situations, complex judgment

Decide with **reversibility** (can you undo it?), **stakes** (money, reputation, legal, harm), **accountability** (who answers for it keeps the decision). Keep a human approval wherever an action is hard to reverse or external. Ask whether the step needs to exist before automating it.

AUGMENT OR REDESIGN

Augment (start here): same workflow, faster steps. **Redesign**: steps merge, move, or disappear, once augmenting works, the process is the bottleneck, and Claude has proven reliable.

1. Pilot one workflow with a small group.
2. Validate on a known answer (redo a past task).
3. Measure against success criteria.
4. Release autonomy gradually.
5. Turn repeated explanations into Projects, instructions, Skills.

COMMUNICATING VALUE AND LIMITS

Bounded task	"Claude drafts the weekly ticket summary."
Benefit, labeled	"We estimate ~2 h/week. We'll measure it."
Limitation	"It can miscategorize. We spot-check 10."
Human owner	"Dana approves before sharing."
Constraint	"No new spend. Changes go through IT."

Avoid "fully automated", "error-free", unmeasured savings. Usage numbers are diagnostics; outcomes prove value.

KNOW WHEN TO ESCALATE

System integrations, pipelines at volume, API, custom connectors or agents	Claude Developer or Architect
Legal, medical, financial, employment judgments	Qualified human professional
Personal or regulated data, new data sources	Legal, privacy, data owner
New tools, integrations, policy exceptions	IT, security, policy owner

D4 TRAPS

- "Claude can handle the whole process, start building" *blind trust*
- "No, it involves customer data" *over-reaction*
- Full autonomy on day one *blind trust*
- Roll out to every team at once *over-engineering*
- Automate a report nobody reads *irrelevant fix*
- Escalate a simple drafting workflow to an Architect *over-engineering*
- Adoption numbers as proof of value *fake safeguard*
- Promise "fully automated, no errors" *blind trust*

D5 Configuration and Knowledge Management

DESCRIPTION

12% · ~7 items · Lesson 05

OBJECTIVES (VERBATIM)

1. Configure Claude Projects with instructions and knowledge sources
2. Manage uploaded knowledge and connectors (e.g., Google Drive, Gmail)

3. Create effective system-level instructions
4. Inform, maintain, and update Claude configurations, knowledge sources, and instructions

WHERE EACH PIECE OF CONTEXT BELONGS

You want Claude to...	Put it in	Scope
Follow house rules in every chat	Organization instructions (admin)	Whole org
Know you, your shorthand, preferences	Personal instructions	All your chats
Behave a set way for one workstream	Project instructions	Every chat in Project
Know reference material for one workstream	Project knowledge	Every chat in Project
Follow a repeatable procedure	Skill	Loads when request matches
Read live or changing files and email	Connector	What your account sees
Recall past conversations	Memory	Per Project or global
Handle today's one-off details	The chat prompt	That message

Instructions say how to behave. **Knowledge** says what to know. A **Skill** says how to do a repeatable task. "Projects store knowledge, skills perform tasks."

PROJECT FACTS THAT SHOW UP IN ITEMS

- Setup: name and **description** · visibility · instructions · knowledge.
- **Claude does not see the description.** A rule typed there has no effect.
- **Context isolation:** a file uploaded in one chat stays in that chat. Add it to knowledge for every chat.
- **Retrieval mode:** near the context limit, Claude searches Project files and pulls relevant parts. Up to 10x capacity on paid plans, with an on-screen indicator.
- Name files descriptively: "Travel-Policy-2026-10.pdf".
- Team/Enterprise sharing: *Can view* (chat, use knowledge), *Can edit* (change instructions and knowledge, manage members), *Owner*.

EFFECTIVE SYSTEM-LEVEL INSTRUCTIONS

Role and context	"You draft replies for Acme's Tier 1 support agents."
Source rule	"Use only the attached help articles. If the answer isn't there, say so."
Process	"Identify the question, find the article, then draft."
Format	"Under 150 words. End with one next step."
Conflict priority	"If brevity and completeness conflict, keep every required fact and cut background."

Use short labeled sections. Keep it concise: every chat carries it.

CONNECTORS (BUILT ON MCP)

- **Claude sees what you see:** your account's permissions.
- Scoped at setup and revocable any time.
- Gmail send/reply/forward and Drive share/move/trash ask for approval by default.
- **Google Drive works in private Projects only.**
- Team/Enterprise admins enable connectors first. Common pattern: read-only first, write access later with risk-owner sign-off.
- Google Docs sync to the latest version. Uploaded files do not update.

Test: "Can you access my Google Drive? List three files in the Marketing folder."

SKILLS AND MEMORY

Skills are folders of instructions and resources that load when a request matches the Skill's description. They need code execution and file creation on. Create one by describing the task to Claude. Install custom Skills only from trusted sources.

Memory helps continuity. It is not a source of truth. Keep authoritative facts in knowledge files.

MAINTAINING CONFIGURATIONS

- **Replace** superseded documents. Old + new side by side gives mixed answers.
- Keep 5–10 test questions with known answers. Rerun after every change.
- Change one thing at a time.
- Give each shared Project an owner and a review rhythm.
- Turn repeated corrections into instructions.

D5 TRAPS

- Put the rule in the Project description *irrelevant fix*
- Add the new policy alongside the old *fake safeguard*
- Tell users to remind Claude which version to use *irrelevant fix*
- Upload in your chat so the team can use it *irrelevant fix*
- Connect the whole Drive, tell Claude to ignore HR *fake safeguard*
- Store a procedure as a knowledge file *irrelevant fix*
- Rebuild the Project from scratch *over-reaction*
- Switch to the most capable model *over-engineering*

D6 Governance, Risk, and Responsible Use

DILIGENCE

15% · ~9 items · Lesson 06

OBJECTIVES (VERBATIM)

1. Identify appropriate and inappropriate use cases
2. Apply data sensitivity, regulatory, and privacy considerations
3. Follow organizational AI policies and governance standards
4. Understand the ethical implications of AI usage

Rule of thumb: make the task safe, then do it. The exam penalizes going ahead unsafely, adding a fake safeguard, and abandoning the task. Sample 3: remove or anonymize customer names and account numbers first, then run the analysis.

APPROPRIATE AND INAPPROPRIATE USES

Claude can handle	Standard replies from documented info, reformatting, summarizing your notes, brainstorming. Delegate, then spot-check.
Claude assists, human decides	Customer letters, survey analysis, a hiring rubric. A person reviews and owns it.
Human handles	Decisions about a person's job, pay, health, credit, or legal standing; sensitive conversations.

High-risk (Usage Policy): legal, healthcare, insurance, finance and lending, employment and housing, academic testing and admissions, media and journalism. Requires a qualified professional in the loop and disclosure of AI involvement. Consumer chatbots must tell users they are talking to AI.

DATA SENSITIVITY: BEFORE YOU UPLOAD

1. What kind of data? Public, internal, confidential, regulated personal data.
 2. What does org policy allow for this data class and tool?
 3. What does the task need? Usually patterns. Rarely names.
- | | |
|------------------------|--|
| Remove | Drop unneeded columns (name, email) |
| Replace | AC-40021 → C001. Keep the key on your machine. |
| Generalize | 2025-02-11 → 2025-02 |
| Scrub free text | Comments hide names, order numbers, places |
| Rare combos | "Only Enterprise customer in LATAM" still identifies |

FAKE SAFEGUARDS

- "Upload it and tell Claude not to retain it." The data already left your hands.
- "Use an Incognito chat." Org retention still applies, and the data was shared.
- "It's internal" or "Legal asked for it." Purpose and seniority do not override policy.
- "Add a disclaimer." Wrong or non-compliant content stays wrong.

PLAN DATA HANDLING (ADMINS SET IT, YOU FOLLOW IT)

Plan	Training on chats	Notes
Free, Pro, Max (consumer)	You choose in Privacy Settings	Allow: kept up to 5 years. Opt out: 30-day standard retention. Paying does not change the choice.
Team, Enterprise, Education, API (commercial)	Not used by default	Enterprise custom retention (30-day min). Admins control memory and connectors.
Incognito (all)	Same as plan	Not in history or memory. Org retention still applies.

Terms change. Verify on Anthropic's privacy pages.

FOLLOWING ORG AI POLICY

- Use the approved workspace for work data, even when a personal account looks faster.
- Missing something? Ask the policy owner. Don't work around it.
- Claude declines and explains why: treat it as information. Don't rephrase to slip past it.
- See a policy-practice gap? Report it and help make the approved path easier.

ETHICS

- **Bias:** coded job-ad language ("digital native"), one-sided framing. Ask Claude to review, then review yourself.
- **Fabrication:** never invent testimonials, reviews, quotes, or data.
- **Accountability:** whoever sends or files the work owns it.
- **Disclosure:** when it changes how people judge the work, when required, and always in high-risk cases. Diligence statement: what AI helped with · which tool · what you reviewed · what you changed · who is responsible.
- Constitutional AI training reduces risk. It does not replace your judgment.

D6 TRAPS

- Upload as-is, the analysis is internal *blind trust*
- Upload and tell Claude not to retain *fake safeguard*
- Skip the analysis / ban AI for the team *over-reaction*
- Let Claude make the hiring decision *blind trust*
- Anonymize, then let Claude decide *fake safeguard*
- Personal Pro account with training off *fake safeguard*
- Escalate every request to Legal *over-engineering*

D7 Troubleshooting and Optimization

DISCERNMENT

10% · ~6 items · Lesson 07

OBJECTIVES (VERBATIM)

1. Identify, diagnose, and resolve issues with underperforming prompts or poor outputs
2. Adjust approach based on feedback and results
3. Optimize workflows for efficiency and effectiveness

TROUBLESHOOTING CHECKLIST: RUN IN ORDER

- 1 **What changed?** Knowledge files, instructions, model or effort, connector, input file, person running it.
- 2 **Context?** Source attached? Fact buried in a long chat? Drift? Context a teammate supplies that you don't?
- 3 **Instructions?** Goal stated? Concrete limits? Conflicting rules?
- 4 **Model and settings?** Light model on hard reasoning, or top model at max effort on routine work?
- 5 **Source of facts?** Attached or connected source, or general knowledge?

Change one thing at a time. Rerun the same 5–10 test inputs with known-good outputs after each change.

ADJUST ON FEEDBACK

- Turn "too long" into "cut the background, keep the three figures".
- Steer early. A quick correction beats a full regenerate.
- Restart when a chat has drifted far.
- Weekly corrections → instructions. Repeated explanations → a Skill.
- The model does not learn from you.

FAILURE → FIRST FIX

Symptom	Likely cause	First fix
Generic output	Missing context	Role, audience, purpose, source, example
Too long / short	Length guessed	State a number
Wrong format	Vague format	Show a filled example
Confident wrong facts	No source	Attach source, ask for quotes, verify
Wrong tone	Default voice	Describe tone, give a sample
Half the ask ignored	Too many asks	Decompose, check each step
Rules forgotten late	Context drift	Summarize, new chat, rules into Project
Essay, no decision	Never asked	Ask for one recommendation
Literal but useless	Letter of the rule	Restate the goal
Agrees with your error	Sycophancy	"Tell me where I'm wrong"
Math errors	LLMs and exact math	Code execution or spreadsheet; check inputs
Works for one colleague	Unstated context	Compare prompts, files, settings; share a Project
Different every run	Underspecified task	Tighten instructions and format; verify at source
Project worse after new files	Stale or conflicting knowledge	Remove superseded files, rerun tests
Missed steps, unfinished	Effort too low	Raise effort
Wordy, scope creep	Effort too high	Lower effort for routine work

OPTIMIZATION LEVERS: LESS TIME AND COST, SAME QUALITY

Model tier	Routine volume on a faster tier; hard reasoning on the top tier
Effort	Lower for routine, higher for complex analysis
Fresh chats	Each turn re-sends everything; new task, new chat
Projects	Stop re-pasting background
Templates, Skills	Save recurring prompts; a Skill loads only when needed
Fewer tools	Turn off unused connectors and features
Real bottleneck	Often a human handoff or approval
Question the step	Confirm someone uses the report first

D7 TRAPS

- Regenerate until it works *irrelevant fix*
 - Switch to the most capable model *over-engineering*
 - Rewrite the whole prompt from scratch *over-reaction*
 - Repeat the instruction in capitals *irrelevant fix*
 - Change model, instructions, and files together *over-engineering*
 - Ask Claude to double-check itself and trust it *fake safeguard*
 - Stop using Claude for this task *over-reaction*
 - Top model at max effort for everything *over-engineering*
- Typical stem: "this used to work and now it doesn't". The best **first** step is to find what changed.

LAB Practical summary: demos and use casesFiles in **assets/demo-files/** · all fictional · run in claude.ai**DEMOS**

Lesson	Demo · files	Key prompt idea · what to look for
00	Walk the 3 official samples	Find the deciding phrase and name each wrong option's pattern. Then ask Claude for 3 new scenarios (1 proportionate answer, 3 distractors) and let it grade your labels.
01 A	Feature picker + two tiers meeting-notes-raw.md	Pick a feature for 8 scenarios. Run "5 bullets + flag the one item a human must check" on Haiku and Sonnet. Judge substance: the CSV with customer emails, the auto-send request.
01 B	Drift → summarize → fresh chat meeting-notes-raw.md	Set 2 rules, add a tangent, ask for a handoff summary under 150 words (goal, decisions, open questions, rules). Check it, start a new chat. Recurring rules go in Project instructions.
02	Raw notes → decomposed prompt meeting-notes-raw.md	Weak baseline, then a tagged extraction prompt with an example row and "Unassigned" / "Not stated" rules. Checkpoint against the notes, specific feedback, email from the checked table, self-review.
03	Spot the errors, verify fictional-regulation-excerpt.md, compliance-summary-draft.md	Claim-by-claim verdict + exact quote + section, regulation only (keep the answer key file claude-summary-with-errors.md out of the chat). Ctrl+F the quotes. Six issues: scope, 12 vs 24 months, two omissions, fake Section 8.2, "start at" vs "up to", unsupported generalization. Sycophancy probe, then an artifact and a table.
04	Map a workflow, brief a stakeholder meeting-notes-raw.md	Claude interviews you (≤5 questions), then builds a delegation map. Legal check and email removal come before upload. Pressure-test auto-send. Note to Tom under 180 words: estimate, limit, approver.
05	Build the Northwind Project project-knowledge-brand-voice.md	Instructions: source rule, [CHECK] placeholders, process, output line. Test happy path, guardrails, chat isolation. Update by replacing the file (12,000+ → 15,000+). A description rule does nothing.
06 A	Anonymize, then analyze churn customer-accounts(-anonymized).csv	Plan redaction from the header + a made-up row. Anonymize in a spreadsheet, upload the clean file. Expect 6/16 churned: Basic 4/6, Pro 2/6, Enterprise 0/4. Count one yourself. Draft a diligence statement.
06 B	High-risk triage	Six cases. High-risk: discharge summary (healthcare), applicant shortlist (employment), rental decline (housing), news article (media). Confirm at the Usage Policy.
07 A	Diagnose a bad prompt project-knowledge-brand-voice.md	"write a linkedin post about our courses" → retry changes nothing → 5-question checklist → rewrite with the source file, length, structure, banned words, and a facts-used list.
07 B	Stale file in the Project brand-voice-2023-edition.md	Add the 2023 guide beside the current one and old prices or partner names leak in. Remove that one file, change nothing else, rerun.
08	Self-assessment	Score 30 objectives 0–3, compute priorities, have Claude draft a 4-week plan, recompute one priority yourself.

USE CASES BY ROLE: SUGGESTED APPROACH**D1 Teacher, 30 essays:** rubric + a model example; review each draft.**D1 Marketer, webinar titles:** 20 titles, then narrow by criteria.**D1 PM, 200 survey responses:** categorize, then count; spot-check counts.**D1 Consultant, market scan:** Research with goals, constraints, sources.**D2 Ops lead, contract clause 14.2:** open the clause yourself; legal reviews.**D2 Marketer, "73% of team leads":** publish only traceable stats.**D2 Consultant, subtotal mismatch:** code execution; check the inputs.**D2 PM, one update, 3 audiences:** exec, team, customer versions; account owner reviews the customer one.**D3 Marketer, 300 tag descriptions:** Haiku or low-effort Sonnet, one tested prompt, spot checks.**D3 Consultant, merger memo:** Opus or Sonnet + extended thinking.**D3 PM, competitor pricing:** Research; open the cited pages.**D3 Ops lead, 70-turn chat drifts:** handoff summary, fresh chat, template into Project instructions.**D4 Ops lead, onboarding emails:** augment: template in a Project, Claude drafts, a person sends, measure 4 weeks.**D4 Marketer, social calendar:** Claude drafts; a person approves; pilot one channel.**D4 Teacher, department feedback:** pilot one teacher, validate on graded reports; grades stay human.**D4 Consultant, client onboarding:** integrations to an Architect; an owner per step.**D5 HR, policy Q&A:** Project with a dated handbook, quoted sections, 10 test questions, old versions replaced.**D5 CS lead, account hub:** private Project + scoped Drive connector (no Drive in shared Projects).**D5 Designer, quiz layout:** Skill = the how; knowledge = the what.**D5 Five writers, inconsistent drafts:** one shared Project, one guide, one owner.**D6 HR, 400 survey comments:** check policy, remove names and team IDs, scrub free text.**D6 Agency, no testimonials:** never invent them; ask real customers.**D6 Finance, duplicate payments:** code vendor names and bank numbers; confirm hits in the finance system.**D6 Instructor, 60 essays:** academic testing is high-risk; Claude drafts feedback, she grades and discloses.**D7 Comms, releases went stiff:** find the new rule, scope it, rerun 3 releases.**D7 Sales ops, limits by noon:** faster tier or lower effort for routine work.**D7 Teacher, uneven feedback:** rubric + annotated samples in a Project; compare to graded essays.**D7 Analyst, totals off:** code execution or a spreadsheet; check inputs.

TIME BUDGET: THREE PASSES

- Pass 1 · 0–80 min** Answer every item in order. Past ~2 min: best answer, flag, move on. No blanks.
- Pass 2 · 80–110 min** Flagged items. Reread for the constraint. Change only with a named reason.
- Sweep · 110–120 min** No blanks. Each multiple-response item has exactly N picks.
- Checkpoints: item 15 ≈ 20 min · 30 ≈ 40 · 45 ≈ 60 · 60 ≈ 80.

MULTIPLE-RESPONSE TACTIC

- Read the count first. Select exactly that many.
 - Judge each option alone: true, and does it answer this stem?
 - **Treat every item as all-or-nothing.** No official source confirms partial credit.
 - Reject near-duplicates and options true in general that miss the scenario.
- "About a quarter of items are multiple-response" is an unofficial community estimate.

HOW TO READ AN ITEM

1. Read the question line first ("best first step", "Choose TWO").
 2. Find the constraint: privacy, stakes, cost/volume, recurrence, recency.
 3. Predict the answer with the core logic, then eliminate by pattern.
 4. Choose the proportionate step that keeps work moving safely.
- Tie-breakers:** the option a careful professional could defend to the risk owner · fix the cause over a per-chat workaround · distrust "always", "never", "guaranteed", "fully automatic".

REGISTER, ID, RESCHEDULE, RETAKE

- **Register:** Partner Academy (partner email) → pay → schedule with Pearson VUE. Get accommodations approved first (allow 10+ days).
- **ID:** government photo ID; name must match exactly. Fix it ≥24 h ahead via certifications-support@anthropic.com, subject "Name Correction Request".
- **Reschedule/cancel ≥48 h ahead.** Refunds need an email to support. No-show or late = fee lost.
- **Online:** system test, clear desk, no phone, watch, or second monitor, stay in webcam view. Accept the NDA.
- **Retakes:** 14 / 30 / 90 days; max 4 per rolling 12 months; fee each time. **Renewal:** free and on time; lapsed = full exam.

FOUR-WEEK STUDY PLAN (~20 H)

Wk	Focus	Do
1	Blueprint, D3, D1	Exam Guide + samples, baseline rubric. Prep courses 1–2, Lessons 00–02. Practice Q1–7, Q19–24.
2	D2, D4	Courses 3–4, Lessons 03–04, their demos. Q8–18, Q25–32.
3	D5, D6, D7	Courses 5–8, Lessons 05–08, their demos. Q33–50. Re-score the rubric.
4	All	All 50 questions in 100 min. Two weakest domains. Logistics (appointment >48 h out, ID, system test). Traps tables, then rest.

45–60 min weekdays + 90 min one weekend day. Prep path: 8 courses, ~6.5 h. Label each miss: **Knowledge** (reread) · **Reading** (slow down) · **Judgment** (drill patterns).

SELF-ASSESSMENT FORMULA

Score each of the 30 objectives: **0** can't explain · **1** can explain · **2** done it with guidance · **3** do it routinely and can say why wrong answers fail.

Priority = (3 – domain average) × domain weight

D2 at 1.5 → $1.5 \times 21 = 31.5$. D7 at 1.0 → $2 \times 10 = 20$. D2 comes first.
Ready to book: ≥2 on every objective, 3 on most of D2, a steady practice score a week apart. Practice bands (uncalibrated): 45–50 strong · 38–44 close · 30–37 keep studying · <30 restart.

OFFICIAL LINKS

Exam Guide v1.0 (PDF) everpath-course-content.s3-accelerate.amazonaws.com/.../Foundations+Exam+Guide.pdf

Certifications overview anthropic-partners.skilljar.com/page/partner-certifications

Certifications FAQ anthropic-partners.skilljar.com/page/faq-certifications

Certification policies anthropic-partners.skilljar.com/page/policies-certifications

Official prep path (8 courses) anthropic-partners.skilljar.com/path/claude-certified-associate-foundations

Pearson VUE for Anthropic pearsonvue.com/us/en/anthropic.html

AI Fluency: Framework and Foundations academy.claude.com/courses/ai-fluency-framework-foundations

Claude 101 academy.claude.com/courses/claude-101

Usage Policy (high-risk use cases) anthropic.com/legal/aup

Prompting best practices platform.claude.com/docs/en/build-with-claude/prompt-engineering/claude-prompting-best-practices

Reduce hallucinations platform.claude.com/docs/en/test-and-evaluate/strengthen-guardrails/reduce-hallucinations

What are Projects? support.claude.com/en/article/s/9517075

Google Workspace connectors support.claude.com/en/articles/10166901

Memory and Incognito support.claude.com/en/articles/11817273

Usage and length limits support.claude.com/en/articles/11647753

Consumer terms update (training choice) anthropic.com/news/updates-to-our-consumer-terms

Support: certifications-support@anthropic.com (name corrections, refunds) · partner-support@anthropic.com (email domain, discounts)

Practice questions in this course are original and do not come from the live exam. Exam content is under NDA. Verify policies at the official pages before booking, because terms change.