

Visual Study Guide: Quiz Questions with Answers

31 exam-style recall questions with worked examples and reference tables,
organized by exam domain.

1 · COVER

Cover the green answer box.
Answer each question out loud
or on paper.

2 · CHECK

Reveal the answer. Mark any
question you missed or
hesitated on.

3 · DRILL

Use the question-only list on the
last page for a closed-book
pass before exam day.

D1 Prompting and Task Execution



14% of exam · 12 questions

D2 Output Evaluation and Validation



21% of exam · 14 questions

D3 Platform and Model Selection



12% of exam · 3 questions

D6 Governance and Responsible Use



15% of exam · 2 questions

Build prompts from explicit components, decompose multi-part work, diagnose weak output and adapt your strategy to the task type.

Q1 What are the five components of a well-built prompt?**ANSWER**

1. **Role**
2. **Context** (including audience)
3. **Task**
4. **Constraints**
5. **Output format**

FIGURE · WEAK VS. STRONG PROMPT**WEAK PROMPT: EVERYTHING LEFT IMPLICIT**

"Write a summary of our quarterly operations."

Claude returns three plausible paragraphs that could describe almost any company. No audience, figures, format or priorities. The output is not wrong. It is unusable, because the prompt specified almost nothing.

STRONG PROMPT: COMPONENTS MADE EXPLICIT

"You are an operations analyst (*role*). I am preparing a one-page update for our regional director, who cares about throughput and cost, not process detail (*context and audience*). Summarize the attached Q3 operations data (*task*), covering only the three metrics that moved more than 10 percent against target (*constraint*). Format as a short headline followed by three bullet points, each one sentence (*output format*)."

Q2 A communications manager must turn a dense 20-page policy change into a staff announcement, a staff FAQ and a short executive briefing. Decompose this into an ordered sequence of steps to run with Claude.**ANSWER**

1. **Extract** the substantive changes and what each one means in practice.
2. **Confirm** the extraction is complete and accurate before building on it.
3. **Draft the staff announcement** from the confirmed change list, tuned to a general audience.
4. **Draft the FAQ**, anticipating the questions staff will ask.
5. **Draft the executive briefing**, compressed to decisions and impact.

Key idea: verify the shared foundation once, then build every deliverable on it.

FIGURE · OUTPUT DEFICIENCIES ARE PROMPT DIAGNOSTICS

Symptom	Likely cause	Fix
Output is generic or off-base	The context was thin	Add the background Claude could not infer
Output answered the wrong question	The task verb was ambiguous	Sharpen the instruction
Output is the wrong length, tone or shape	A constraint or the format was missing	Add it
Output is close but misses on one section	—	Iterate on that section only. Do not discard a draft that is mostly right

Q3 Claude's output is generic and off-base. What is the likely cause and the fix?

ANSWER Cause: **thin context**. Fix: **add the background Claude could not infer**.

Source: Prompt diagnostics

Q4 Claude answered the wrong question. What is the likely cause and the fix?

ANSWER Cause: **an ambiguous task verb**. Fix: **sharpen the instruction**.

Source: Prompt diagnostics

Q5 A draft is close but misses on one section. What should you do?

ANSWER **Iterate on that section only**. Do not discard a draft that is mostly right.

Source: Prompt diagnostics

WORKED EXAMPLE · A LIVE ITERATION CYCLE

ROUND 1 PROMPT

"Write a follow-up email to the client about the delayed deliverable."

ROUND 1 OUTPUT

A generic, slightly defensive three-paragraph email that never says when the deliverable will arrive or why it slipped. **Diagnosis:** thin context (no reason, no new date) and no tone constraint.

ROUND 2 PROMPT

"Write a follow-up to the client about the two-day delay on the analytics deliverable. The cause was a data-quality issue we have now fixed; new delivery is Thursday. Tone: accountable, not over-apologetic. Keep it under 120 words."

ROUND 2 OUTPUT

A tight, accountable note with the cause, the new date and a confident close. **Diagnosis:** strong. Only the subject line is missing.

ROUND 3 PROMPT

"Good. Add a subject line that signals resolution, not just delay."

The lesson: fix the specific component the output exposed. Do not restart from scratch.

Q6 What are the four core task types for everyday chat use of Claude?

ANSWER Analysis · Research · Drafting · Brainstorming

Q7 When you set up an analysis prompt, what three things do you spell out?

ANSWER What to measure, the standard to measure it against and how to handle ambiguity. Low creative latitude, high specification.

Source: Prompt configuration by task type

Q8 When you set up a research prompt, what do you define up front and what do you ask for in the output?

ANSWER Define the question, its boundaries and whether current sources are required. Ask for citations so you can check the claims.

Source: Prompt configuration by task type

Q9 In a drafting prompt, what do you specify and what do you leave to Claude?

ANSWER You specify audience, tone and format. Claude finds the phrasing. You control the shape; Claude fills it.

Source: Prompt configuration by task type

Q10 How do you set up a brainstorming chat, and why keep it loose?

ANSWER Give the goal and the boundaries, then ask for many varied ideas before narrowing. Too many constraints kill the variety you want.

Source: Prompt configuration by task type

FIGURE · STRATEGY QUICK REFERENCE

Task type	What to tighten	What to loosen	Latitude
Analysis	Criteria, standards, scope	Phrasing	Low
Research	Question, sources, citations	Synthesis approach	Low–med
Drafting	Audience, tone, format	Word choice	Medium
Brainstorming	Goal and guardrails only	Quantity and direction	High

Research tip: web search in chat covers quick currency needs. Deep multi-source investigation points to the Research feature (paid plans).

Q11 Name the three ways to format Claude's output by purpose.

ANSWER Inline · Artifact · Structured format (table, JSON, CSV and similar)

Source: Format by purpose

Q12 What separates an artifact from an inline reply?

ANSWER An artifact is a **standalone deliverable** you edit and reuse. An inline reply is contextual and gets used up in the chat.

Source: Format by purpose

D2 Output Evaluation and Validation

Exam weight **21%**

Spot hallucination patterns, tune prompts for verifiability, check output with concrete techniques and decide what must never ship without human review.

Q13 Describe in one sentence the hallucination pattern called *plausible-but-unsupported claims*.

ANSWER A statement that sounds reasonable and fits the topic but has **no basis in the source or in fact**. It is the most dangerous kind because nothing about it looks wrong.

Q14 Your prompt implies a preferred answer and Claude agrees a little too readily on a question that should be open. What is this an example of?

ANSWER **Confirmation bias in framing**: the way you phrased the question steered the answer.

Q15 What are the three techniques to tune a prompt for verifiability, and how do you implement each?

ANSWER

1. **Permit "I don't know"**: tell Claude it may say so (works well in a system prompt or Project instructions).
2. **Restrict to provided sources**: attach the documents or links and tell Claude to answer only from them.
3. **Require auditable citations**: ask for direct quotes from the PDF or traceable links to the exact section of a page.

Q16 Name three ways to check Claude's output.

ANSWER **Quote first, then analyze** · **Best-of-N comparison** · **Validate against an authoritative source**

Source: Anthropic docs: Reduce hallucinations

Q17 Before Claude analyzes a long document, what should you ask it to do first?

ANSWER Extract the **word-for-word quotes** that support the task, then analyze using only those quotes. Anthropic suggests this for documents over ~20k tokens.

Source: Anthropic docs: Reduce hallucinations

Q18 Why does quote-first analysis make Claude's output easier to check?

ANSWER Each conclusion ties to a pulled quote, so the reasoning and the errors become **visible**. You can check each quote against the source.

Source: Anthropic docs: Reduce hallucinations

Q19 Best-of-N: after re-running the same request several times, how do you read the results?

ANSWER **Agreement** raises confidence. **Divergence** marks a soft spot that needs a human look (a possible hallucination).

Source: Anthropic docs: Reduce hallucinations

Q20 Why isn't a second Claude response enough to validate a claim that matters?

ANSWER It is **not independent** and can repeat the same error. Check against a **trusted external reference**.

Source: Anthropic docs: "Always validate critical information"

Q21 You're about to act on a figure Claude produced in a spreadsheet. Which in-product aid helps you validate it?

ANSWER **Claude for Excel's cell-level citations**, which tie each figure back to its input cells.

Q22 What belongs on the *do-not-ship-without-review* list? (Decide it in advance and treat it as fixed.)

ANSWER

- Final client deliverables
- Audit-critical or financially material calculations
- Anything with regulated, confidential or highly sensitive data
- Public or legal communications where a misstatement has lasting consequences

FIGURE · FOUR RISK THRESHOLDS FOR REVIEWING AI OUTPUT

Threshold	What to ask
Stakes	What is the cost if this is wrong? High-cost errors demand human review regardless of how confident the output appears.
Reversibility	Can the action be undone? An irreversible step (a sent client deliverable, a filed report) clears a higher bar than a draft you can revise.
Audience	Who sees it? External, executive and regulatory audiences raise the review requirement above internal working drafts.
Regulatory exposure	Does a rule, contract or law govern this? Regulated content carries obligations that AI assistance does not remove.

Q23 Risk threshold Stakes: what question do you ask?

ANSWER **What is the cost if this is wrong?** High-cost errors demand human review however confident the output looks.

Source: Four risk thresholds

Q24 Risk threshold Reversibility: what question do you ask?

ANSWER **Can the action be undone?** A sent client deliverable or a filed report clears a higher bar than a revisable draft.

Source: Four risk thresholds

Q25 Risk threshold Audience: what question do you ask?

ANSWER **Who sees it?** External, executive and regulatory audiences raise the review requirement above internal drafts.

Source: Four risk thresholds

Q26 Risk threshold Regulatory exposure: what question do you ask?

ANSWER **Does a rule, contract or law govern this?** Regulated content carries obligations that AI assistance does not remove.

Source: Four risk thresholds

D3 Platform and Model Selection

Exam weight **12%**

Know what a context window holds and how the chat product's limits differ from the model's.

FIGURE · WHAT FILLS ONE REQUEST'S CONTEXT WINDOW

System prompt

Messages · files · tool results

Tool definitions

Output

Thinking

input + output + reasoning ≤ model window. The chat product (claude.ai, ChatGPT, Gemini app) may cap you lower than the model's advertised maximum.

Q27 What counts toward a model's context window in a single API request?

ANSWER Everything sent **plus** everything generated: system prompt + messages (including tool results and files) + tool definitions + output + thinking tokens.

Source: Anthropic docs: Context windows

Q28 Model context window vs. chat-product context window: which one limits you in the app, and how do they relate?

ANSWER The **product's cap** limits you. The model window is the hard ceiling for one request. The product sets its own cap per plan and model, and that cap can sit **below** the model's maximum.

Source: Anthropic docs; Claude Help Center

Q29 A chat app loses early details in a long conversation even though its model advertises 1M tokens. What should you suspect first?

ANSWER The **product's cap** for your plan and model, plus the app **trimming old turns** (rolling first-in-first-out). The model's advertised max is not what the app gives you.

Source: Anthropic docs: Context windows

Match data-handling controls to the sensitivity of the task.

Q30 When is redaction the wrong tool for making data usable?

ANSWER When the task **genuinely depends on the sensitive specifics** you would remove.

Source: Data sensitivity and privacy controls

Q31 Incognito keeps a chat out of Memory and chat history. Which separate control still applies to Incognito chats and can surface them in org data exports?

ANSWER **The organization's data-retention policy.**

Source: Data sensitivity and privacy controls

Closed-book drill

Questions only · answers in the domain sections

1. **D1** What are the five components of a well-built prompt?
2. **D1** A communications manager must turn a dense 20-page policy change into a staff announcement, a staff FAQ and a short executive briefing. Decompose this into an ordered sequence of steps to run with Claude.
3. **D1** Claude's output is generic and off-base. What is the likely cause and the fix?
4. **D1** Claude answered the wrong question. What is the likely cause and the fix?
5. **D1** A draft is close but misses on one section. What should you do?
6. **D1** What are the four core task types for everyday chat use of Claude?
7. **D1** When you set up an analysis prompt, what three things do you spell out?
8. **D1** When you set up a research prompt, what do you define up front and what do you ask for in the output?
9. **D1** In a drafting prompt, what do you specify and what do you leave to Claude?
10. **D1** How do you set up a brainstorming chat, and why keep it loose?
11. **D1** Name the three ways to format Claude's output by purpose.
12. **D1** What separates an artifact from an inline reply?
13. **D2** Describe in one sentence the hallucination pattern called *plausible-but-unsupported claims*.
14. **D2** Your prompt implies a preferred answer and Claude agrees a little too readily on a question that should be open. What is this an example of?
15. **D2** What are the three techniques to tune a prompt for verifiability, and how do you implement each?
16. **D2** Name three ways to check Claude's output.
17. **D2** Before Claude analyzes a long document, what should you ask it to do first?
18. **D2** Why does quote-first analysis make Claude's output easier to check?
19. **D2** Best-of-N: after re-running the same request several times, how do you read the results?
20. **D2** Why isn't a second Claude response enough to validate a claim that matters?
21. **D2** You're about to act on a figure Claude produced in a spreadsheet. Which in-product aid helps you validate it?
22. **D2** What belongs on the *do-not-ship-without-review* list? (Decide it in advance and treat it as fixed.)
23. **D2** Risk threshold Stakes: what question do you ask?
24. **D2** Risk threshold Reversibility: what question do you ask?
25. **D2** Risk threshold Audience: what question do you ask?
26. **D2** Risk threshold Regulatory exposure: what question do you ask?
27. **D3** What counts toward a model's context window in a single API request?
28. **D3** Model context window vs. chat-product context window: which one limits you in the app, and how do they relate?
29. **D3** A chat app loses early details in a long conversation even though its model advertises 1M tokens. What should you suspect first?
30. **D6** When is redaction the wrong tool for making data usable?
31. **D6** Incognito keeps a chat out of Memory and chat history. Which separate control still applies to Incognito chats and can surface them in org data exports?